

GPU→CPU Transfer Throughput vs Document Length
(varying max_num_batched_tokens)

